

Sequential Algorithms

Classical Algorithm Design:

- One machine/CPU/process/... doing a computation

RAM (Random Access Machine):

- Basic standard model
- Unit cost basic operations
- Unit cost access to all memory cells

Sequential Algorithm / Program:

- Sequence of operations
(executed one after the other)

Parallel and Distributed Algorithms

Today's computers/systems are not sequential:

- Even cell phones have several cores
- Future systems will be highly parallel on many levels
- This also requires appropriate algorithmic techniques

Goals, Scenarios, Challenges:

- Exploit parallelism to speed up computations
- Shared resources such as memory, bandwidth, ...
- Increase reliability by adding redundancy
- Solve tasks in inherently decentralized environments
- ...

Parallel and Distributed Systems

- Many different forms
- Processors/computers/machines/... communicate and share data through
 - Shared memory or message passing
- Computation and communication can be
 - Synchronous or asynchronous
- Many possible **topologies** for message passing
- Depending on system, various types of faults

Challenges

Algorithmic and theoretical challenges:


- How to parallelize computations
- Scheduling (which machine does what)
- Load balancing
- Fault tolerance
- Coordination / consistency
- Decentralized state
- Asynchrony
- Bounded bandwidth / properties of comm. channels
- ...

Models

- A large variety of models, e.g.:
- **PRAM** (Parallel Random Access Machine)
 - Classical model for parallel computations
- **Shared Memory**
 - Classical model to study coordination / agreement problems, distributed data structures, ...
- **Message Passing** (fully connected topology)
 - Closely related to shared memory models
- Message Passing in **Networks**
 - Decentralized computations, large parallel machines, comes in various flavors...

PRAM

- Parallel version of RAM model
- p processors, shared random access memory




- Basic operations / access to shared memory cost 1
- Processor operations are synchronized
- Focus on parallelizing computation rather than cost of communication, locality, faults, asynchrony, ...

Parallel Computations


Sequential Computation:

- Sequence of operations



Parallel Computation:

- Directed Acyclic Graph (DAG)




Parallel Computations

$\underline{T_p}$: time to perform comp. with p procs

- $\underline{T_1}$: **work** (total # operations)
 - Time when doing the computation sequentially
- $\underline{T_\infty}$: **critical path / span**
 - Time when parallelizing as much as possible
- **Lower Bounds:**

$$\underline{T_p} \geq \frac{\underline{T_1}}{p}, \quad \underline{T_p} \geq \underline{T_\infty}$$



Parallel Computations

T_p : time to perform comp. with p procs


- Lower Bounds:

$$T_p \geq \frac{T_1}{p}, \quad T_p \geq T_\infty$$

- Parallelism: $\frac{T_1}{T_\infty}$
 - maximum possible speed-up

- Linear Speed-up:

$$\frac{T_p}{T_1} = \Theta(p)$$



Scheduling

- How to assign operations to processors?
- Generally an online problem
 - When scheduling some jobs/operations, we do not know how the computation evolves over time

Greedy (offline) scheduling:

- Order jobs/operations as they would be scheduled optimally with ∞ processors (topological sort of DAG)
 - Easy to determine: With ∞ processors, one always schedules all jobs/ops that can be scheduled
- Always schedule as many jobs/ops as possible
- Schedule jobs/ops in the same order as with ∞ processors
 - i.e., jobs that become available earlier have priority

Brent's Theorem

Brent's Theorem: On p processors, a parallel computation can be performed in time

$$T_p \leq \frac{T_1 - T_\infty}{p} + T_\infty.$$

Proof:

- Greedy scheduling achieves this...
- #operations scheduled with ∞ processors in round i : x_i

$$\sum_{i=1}^{T_\infty} x_i = T_1$$

done skip x_i

time steps for greedy to do x_i jobs

$$\left\lceil \frac{x_i}{p} \right\rceil \leq \frac{x_i}{p} + \frac{p-1}{p} = 1 + \frac{x_i-1}{p}$$

$$T_p \leq \sum_{i=1}^{T_\infty} \left(1 + \frac{x_i-1}{p} \right) = T_\infty + \underbrace{\sum_{i=1}^{T_\infty} \frac{x_i}{p}}_{T_1/p} - \underbrace{\frac{1}{p} \sum_{i=1}^{T_\infty} 1}_{T_\infty/p} = T_\infty + \frac{T_1 - T_\infty}{p}$$

Brent's Theorem

Brent's Theorem: On p processors, a parallel computation can be performed in time

$$T_p \leq \frac{T_1 - T_\infty}{p} + T_\infty.$$

Proof:

- Greedy scheduling achieves this...
- #operations scheduled with ∞ processors in round i : x_i

Brent's Theorem

Brent's Theorem: On p processors, a parallel computation can be performed in time

$$T_p \leq \frac{T_1 - T_\infty}{p} + T_\infty.$$

Corollary: Greedy is a 2-approximation algorithm for scheduling.

$$T_p \geq \frac{T_1}{p} \quad T_p \geq T_\infty \quad T_p < \frac{T_1}{p} + T_\infty \leq 2 \cdot \max \left\{ \frac{T_1}{p}, T_\infty \right\}$$

$T_p \geq \dots$

Corollary: As long as the number of processors $p = O(T_1/T_\infty)$, it is possible to achieve a linear speed-up.

$$\frac{T_1}{p} = \Omega(T_\infty)$$

PRAM

Back to the PRAM:

- Shared random access memory, synchronous computation steps
- The PRAM model comes in variants...

EREW (exclusive read, exclusive write): *to the same memory cell*

- Concurrent memory access by multiple processors is not allowed
- If two or more processors try to read from or write to the same memory cell concurrently, the behavior is not specified

CREW (concurrent read, exclusive write):

- Reading the same memory cell concurrently is OK
- Two concurrent writes to the same cell lead to unspecified behavior
- This is the first variant that was considered (already in the 70s)

PRAM

The PRAM model comes in variants...

CRCW (concurrent read, concurrent write):

- Concurrent reads and writes are both OK
- Behavior of concurrent writes has to be specified
 - – Weak CRCW: concurrent write only OK if all processors write 0
 - – Common-mode CRCW: all processors need to write the same value
 - – Arbitrary-winner CRCW: adversary picks one of the values
 - – Priority CRCW: value of processor with highest ID is written
 - – Strong CRCW: largest (or smallest) value is written
- The given models are ordered in strength:
weak \leq common-mode \leq arbitrary-winner \leq priority \leq strong

Some Relations Between PRAM Models

Theorem: A parallel computation that can be performed in time t , using p processors on a strong CRCW machine, can also be performed in time $O(t \log p)$ using p processors on an EREW machine.

- Each (parallel) step on the CRCW machine can be simulated by $O(\log p)$ steps on an EREW machine

Theorem: A parallel computation that can be performed in time t , using p probabilistic processors on a strong CRCW machine, can also be performed in expected time $O(t \log p)$ using $O(p/\log p)$ processors on an arbitrary-winner CRCW machine.

- The same simulation turns out more efficient in this case

Some Relations Between PRAM Models

Theorem: A computation that can be performed in time t , using p processors on a strong CRCW machine, can also be performed in time $O(t)$ using $O(p^2)$ processors on a weak CRCW machine

Proof:

- **Strong:** largest value wins, **weak:** only concurrently writing 0 is OK

$$t = O(1) \quad p = \sqrt{n}$$

$$O(t) = O(1) \quad O(p^2) = O(n)$$

Some Relations Between PRAM Models

Theorem: A computation that can be performed in time t , using \underline{p} processors on a strong CRCW machine, can also be performed in time $O(t)$ using $O(p^2)$ processors on a weak CRCW machine

Proof:

- **Strong:** largest value wins, **weak:** only concurrently writing 0 is OK

Computing the Maximum

Observation: On a strong CRCW machine, the maximum of n values can be computed in $O(1)$ time using n processors

- Each value is concurrently written to the same memory cell

Lemma: On a weak CRCW machine, the maximum of n integers between 1 and \sqrt{n} can be computed in time $O(1)$ using $O(n)$ proc.

Proof:

$\max i \text{ such that } f_i = 0$

- We have \sqrt{n} memory cells $f_1, \dots, f_{\sqrt{n}}$ for the possible values
- Initialize all $f_i := 1$
- For the n values x_1, \dots, x_n , processor j sets $f_{x_j} := 0$
 - Since only zeroes are written, concurrent writes are OK
- Now, $f_i = 0$ iff value i occurs at least once
- Strong CRCW machine: max. value in time $O(1)$ w. $O(\sqrt{n})$ proc.
- Weak CRCW machine: time $O(1)$ using $O(n)$ proc. (prev. lemma)

Computing the Maximum

Theorem: If each value can be represented using $\underline{O(\log n)}$ bits, the maximum of n (integer) values can be computed in time $O(1)$ using $\underline{O(n)}$ processors on a weak CRCW machine.

Proof:

- First look at $\frac{\log_2 n}{2}$ highest order bits
- The maximum value also has the maximum among those bits
- There are only \sqrt{n} possibilities for these bits
- max. of $\frac{\log_2 n}{2}$ highest order bits can be computed in $O(1)$ time
- For those with largest $\frac{\log_2 n}{2}$ highest order bits, continue with next block of $\frac{\log_2 n}{2}$ bits, ...

