

# *Why Election?*

- ❖ **Example 1: Your Bank maintains multiple servers in their cloud, but for each customer, one of the servers is responsible, i.e., is the leader**
  - ❖ What if there are two leaders per customer?
    - ❖ Inconsistency
  - ❖ What if servers disagree about who the leader is?
    - ❖ Inconsistency
  - ❖ What if the leader crashes?
    - ❖ Unavailability

# *Why Election?*

- ❖ **Example 2: In the sequencer-based algorithm for total ordering of multicasts, the "sequencer" = leader**
- ❖ **Example 3: Group of cloud servers replicating a file need to elect one among them as the primary replica that will communicate with the client machines**
- ❖ **Example 4: Group of NTP servers: who is the root server?**

# **What is Election?**

- ❖ In a group of processes, elect a **Leader** to undertake special tasks.
- ❖ What happens when a leader fails (crashes)
  - ❖ Some (at least one) process detects this (how?)
  - ❖ Then what?
- ❖ Focus of this lecture: **Election algorithm**
  1. Elect one leader only among the non-faulty processes
  2. All non-faulty processes agree on who is the leader

# **System Model/Assumptions**

- ❖ Any process can **call** for an **election**.
- ❖ A process can **call** for **at most one election at a time**.
- ❖ Multiple processes can **call** an election **simultaneously**.
  - ❖ All of them together must yield a single leader only
- ❖ The result of an election should not depend on which process calls for it.
- ❖ Messages are eventually delivered.

# **Problem Specification**

- ❖ At the end of the election protocol, the non-faulty process with the best (highest) election attribute value is elected.
  - ❖ Attribute examples: leader has highest id or address. Fastest cpu. Most disk space. Most number of files, etc.
- ❖ Protocol may be initiated anytime or after leader failure
- ❖ A *run* (execution) of the election algorithm must always guarantee at the end:
  - Safety:  $\forall$  non-faulty p: (p's elected = (q: a particular non-faulty process with the best attribute value) or  $\perp$ )
  - Liveness:  $\forall$  election: (election terminates)  
 $\quad \& \forall p:$  non-faulty process, p's elected is not  $\perp$


# Algorithm 1: Ring Election

- ❖ N Processes are organized in a logical ring
  - ❖  $p_i$  has a communication channel to  $p_{(i+1) \bmod N}$
  - ❖ All messages are sent clockwise around the ring.
- ❖ Any process  $p_i$  that discovers the old coordinator has failed initiates an “election” message that contains  $p_i$ ’s own id:attr. This is the *initiator* of the election.
- ❖ When a process  $p_i$  receives an *election* message, it compares the attr in the message with its own attr.
  - ❖ If the arrived attr is greater,  $p_i$  forwards the message.
  - ❖ If the arrived attr is smaller and  $p_i$  has not yet forwarded an election message, it overwrites the message with its own id:attr, and forwards it.
  - ❖ If the arrived id:attr matches that of  $p_i$ , then  $p_i$ ’s attr must be the greatest (why?), and it becomes the new coordinator. This process then sends an “elected” message to its neighbor with its id, announcing the election result.
- ❖ When a process  $p_i$  receives an *elected* message, it
  - ❖ sets its variable *elected*,  $\leftarrow$  id of the message.
  - ❖ forwards the message, unless it is the new coordinator.

# Ring-Based Election: Example

(In this example, attr:=id)

- In the example: The election was started by process 17. The highest process identifier encountered so far is 24. (final leader will be 33)
- The worst-case scenario occurs when the counter-clockwise neighbor (@ the initiator) has the highest attr.




# Ring-Based Election: Analysis

- ❖ The worst-case scenario occurs when the counter-clockwise neighbor has the highest attr.

In a ring of N processes, in the worst case:

- ❖ A total of  $N-1$  messages are required to reach the new coordinator-to-be (election messages).
- ❖ Another  $N$  messages are required until the new coordinator-to-be ensures it is the new coordinator (election messages – no changes).
- ❖ Another  $N$  messages are required to circulate the elected messages.
- ❖ Total Message Complexity =  $3N-1$
- ❖ Turnaround time =  $3N-1$




# **Correctness?**

**Assume – no failures happen during the run of the election algorithm**

- **Safety and Liveness are satisfied.**

**What happens if there are failures during the election run?**

# Example: Ring Election




May not terminate when process failure occurs during the election!  
Consider above example where attr == id

Does not satisfy liveness

## **Algorithm 2: Modified Ring Election**

- ❖ Processes are organized in a logical ring.
- ❖ Any process that discovers the coordinator (leader) has failed initiates an “**election**” message.
- ❖ The message is circulated around the ring, bypassing failed processes.
- ❖ Each process appends (adds) its id:attr to the message as it passes it to the next process (without overwriting what is already in the message)
- ❖ Once the message gets back to the initiator, it elects the process with the best election attribute value.
- ❖ It then sends a “**coordinator**” message with the id of the newly-elected coordinator. Again, each process adds its id to the end of the message, and records the coordinator id locally.
- ❖ Once “**coordinator**” message gets back to initiator,
  - ❖ election is over if would-be-coordinator’s id is in id-list.
  - ❖ else the algorithm is repeated (handles election failure).

# Example: Ring Election



# **Modified Ring Election**

- Supports concurrent elections – an initiator with a lower id blocks other initiators' election messages
- Reconfiguration of ring upon failures
  - Can be done if all processes “know” about all other processes in the system (Membership list! – MP2)
- If initiator non-faulty ...
  - How many messages?  $2N$
  - What is the turnaround time?  $2N$
  - Size of messages?  $O(N)$
- How would you redesign the algorithm to be fault-tolerant to an initiator's failure?
  - One idea: Have the initiator's successor wait a while, timeout, then re-initiate a new election. Do the same for this successor's successor, and so on...
  - What if timeouts are too short... starts to get messy

# **Leader Election Is Hard**

- The Election problem is related to the ***consensus problem***
- Consensus is impossible to solve with 100% guarantee in an asynchronous system with no bounds on message delays and arbitrarily slow processes
- So is leader election in fully asynchronous system model
- Where does the modified Ring election start to give problems with the above asynchronous system assumptions?
  - $p_i$  may just be very slow, but not faulty (yet it is not elected as leader!)
  - Also slow initiator, ring reorganization

# ***Algorithm 3: Bully Algorithm***


- ❖ Assumptions:
  - ❖ Synchronous system
    - ❖ All messages arrive within  $T_{trans}$  units of time.
    - ❖ A reply is dispatched within  $T_{process}$  units of time after the receipt of a message.
    - ❖ if no response is received in  $2T_{trans} + T_{process}$ , the process is assumed to be faulty (crashed).
  - ❖ attr==id
  - ❖ Each process knows all the other processes in the system (and thus their id's)

## **Algorithm 3: Bully Algorithm**


- ❖ When a process finds the coordinator has failed, if it knows its id is the highest, it elects itself as coordinator, then sends a **coordinator** message to all processes with lower identifiers than itself
- ❖ A process initiates election by sending an **election** message to only processes that have a higher id than itself.
  - If no answer within timeout, send coordinator message to lower id processes → Done.
  - if any answer received, then there is some non-faulty higher process → so, wait for coordinator message. If none received after another timeout, start a new election.
- ❖ A process that receives an “election” message replies with **answer** message, & starts its own election protocol (unless it has already done so)

# Example: Bully Election


answer=OK




1. P2 initiates election




2. P2 receives answers




3. P3 & P4 initiate election



4. P3 receives reply




5. P4 receives no reply



5. P4 announces itself

# The Bully Algorithm with Failures

The coordinator  $p_4$  fails and  $p_1$  detects this



# *Analysis of The Bully Algorithm*

- **Best case scenario:** The process with the second highest id notices the failure of the coordinator and elects itself.
  - $N-2$  coordinator messages are sent.
  - Turnaround time is one message transmission time.

# *Analysis of The Bully Algorithm*

- **Worst case scenario: When the process with the lowest id in the system detects the failure.**
  - N-1 processes altogether begin elections, each sending messages to processes with higher ids.
    - » i-th highest id process sends i-1 election messages
  - The message overhead is  $O(N^2)$ .
  - Turnaround time is approximately 5 message transmission times if there are no failures during the run:
    1. Election message from lowest id process
    2. Answer to lowest id process from 2<sup>nd</sup> highest id process
    3. Election from 2nd highest id process
    4. Timeout for answers @ 2nd highest id process
    5. Coordinator message from 2<sup>nd</sup> highest id process

# **Summary**

- Coordination in distributed systems requires a leader process
- Leader process might fail
- Need to (re-) elect leader process
- Three Algorithms
  - Ring algorithm
  - Modified Ring algorithm
  - Bully Algorithm